

High Performance Computing (HPC)


- Introduction to HPC and why it's useful
- How do you parallelize your code?
- The practice of running software on a cluster



Photo: Argonne National Laboratory

Introduction to HPC and why it's useful

Gordon Moore




Born January 3, 1929 (age 84)
San Francisco, California, USA

Nationality American

Alma mater University of California, Berkeley;
California Institute of Technology

Occupation Chairman Emeritus, Intel Corporation

Net worth ▲ \$4 billion USD (2011)^[1]



Introduction to HPC and why it's useful


- Embarrassingly parallel problems
 - Graphics
 - Simulations with multiple parameters
- Non embarrassingly parallel problems
 - Fluid dynamics
 - A lot of the tasks run by a single program




How do you parallelize your code? !

- `as.numeric(Sys.getenv("PBS_ARRAY_INDEX"))`
– Should be used in your code to give a simulation number

Before



Now



Using your PC

```
for ( i in 1 : 10 )
{
  do_simulation( i )
}
```

Shell script on
the cluster


Using HPC

```
i <- as.numeric(Sys.getenv("PBS_ARRAY_INDEX"))
do_simulation( i )
```


How do you parallelize your code?

- `as.numeric(Sys.getenv("PBS_ARRAY_INDEX"))`
– Should be used in your code to give a simulation number

Before




Now




How do you parallelize your code?

- `as.numeric(Sys.getenv("PBS_ARRAY_INDEX"))`
 - Should be used in your code to give a simulation number
- Pseudo random numbers
 - Given a certain random number seed, you get the same sequence of random numbers every time
 - Each simulation should have a different seed




Handling memory



Handling memory

- Save your results in memory and then write to disk at the end



Handling memory




- Save your results in memory and then write to disk at the end
- Output your code to a series of files
- Write local code to read in your series of files automatically



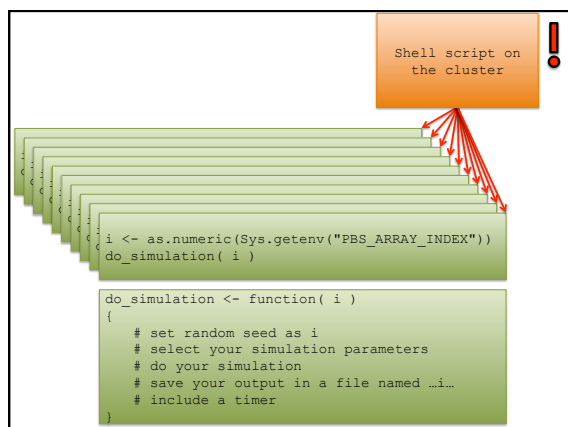
Handling memory

- Save your results in memory and then write to disk at the end
- Output your code to a series of files
- Write local code to read in your series of files automatically

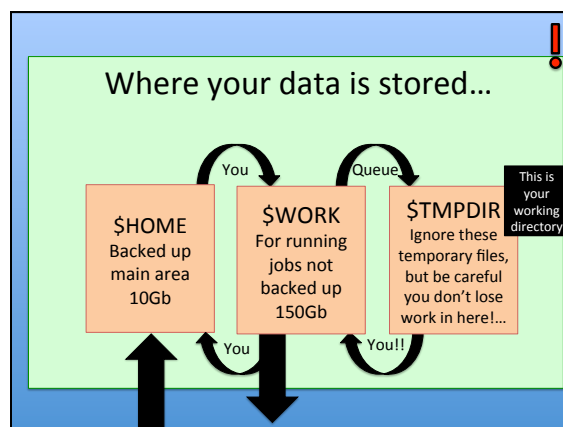
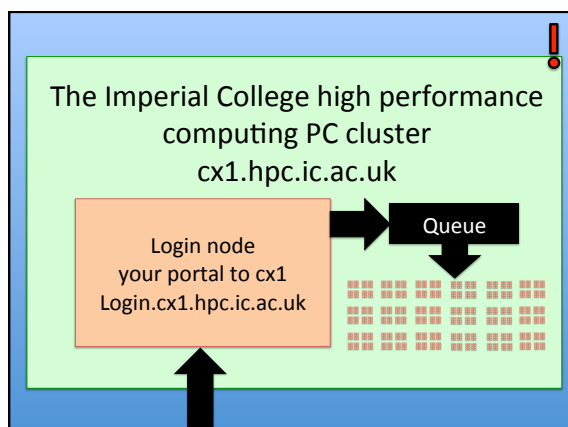



Handling memory

- Save your results in memory and then write to disk at the end
- Output your code to a series of files
- Write local code to read in your series of files automatically
- Build a timer into your code
- Test your code locally to know your memory and time requirements



The practice of running code on a cluster



Step 1: get your code onto the cluster

- Use sftp: from the directory of your code in a shell window type
 - sftp `username@login.cx1.hpc.ic.ac.uk`
 - You will be asked for your standard cluster password
 - cd `$HOME`
 - put `filename.R`
 - exit
- Your software is now on your home directory of the cluster which is backed up regularly

Step 2: log into the cluster

- Use ssh: from a shell window type
 - Ssh `-l username login.cx1.hpc.ic.ac.uk`
 - You will be asked for your standard cluster password
 - Now it's as though you were sitting with a shell open at the login node.
 - ls (will list the files in \$HOME)
 - mkdir `foldername` (make a new folder)
 - mv `filename $HOME/foldername` (move)
 - cd `foldername` (change directory)
 - cat `filename` (see your file to check it's contents)
 - cp `filename $WORK` (make a copy in \$WORK)
 - cd `$WORK` (change directory to \$WORK)
- Your software is now in two copies: in a folder in your home directory of the cluster and in the work area ready for running.

Step 3: make a file for your shell script !

- This is the list of instructions to be executed when you get to the front of the queue is written in shell script. It should be a .sh file
- **Never** run code on the login node – always write a shell script and wait in the queue.
- If you type `cat > filename.sh`
- You will then get the chance to type text (pressing enter for new lines) and the cat command will make the file containing the text that you typed.
- When you are finished typing the contents of your new file press control and D to complete the process.
- Type `cat filename.sh` to check that your file is correct before submitting it to the queue.

Step 3 continued: your shell script file !

```
#!/bin/bash
#PBS -l walltime=12:00:00
#PBS -l select=1:ncpus=1:mem=800mb
module load R/2.13.0
module load intel-suite
echo "R is about to run"
R --vanilla < $WORK/Rtest/ForwardsNTC_V5.R
mv datafilename* $WORK
echo "R has finished running"
# this is a comment at the end of the file
```

You can run Python code too –
just use different commands here

Step 4: submitting your job to the cluster !

- You are now in the \$WORK directory with your code and shell script both written.
- To submit your job type
`qsub -J 1-32 filename.sh`
`qstat` (S changes from Q to B when running)
- If you want to delete a job
`qstat`
`qdel job-id[]` (the [] is for array jobs only)
- `qstat` will give you a list of jobs and you would get the job-id from there.

Step 5: check that all is well !

- Wait 5-10 minutes then check that nothing has gone wrong.
- `qstat` (is your job running still)
- `ls` (are output files as expected)
- `cat filename.sh.ejob-id.index`
(are error files empty?)
- `cat filename.sh.ojob-id.index`
(are standard output files as expected)
- `qstat` (is your job running still)
- `exit` (you're done for now come back later)

Step 6: Getting your results back !

- `qstat` (is your job running still)
- `cd $WORK`
- `ls` (output files as expected?)
- `cat output filename` (contents as expected?)
- `cat filename.sh.ejob-id.index`
(error files empty?)
- `cat filename.sh.ojob-id.index`
(standard output files as expected?)
- `tar czvf filename.tgz *`
- `mv filename.tgz $HOME/foldername`


Step 6 continued: sftp to get results !

- `exit`
- Use sftp: from a new directory on your own computer of where you want the results to be. Open a shell and type ...
 - `sftp username@login.cx1.hpc.ic.ac.uk`
 - You will be asked for your standard cluster password
 - `cd $HOME/foldername`
 - `get filename.tgz`
 - `exit`
- Your results are now all on your own computer
 - `tar xzvf filename.tgz`
- Your results are now complete uncompressed and ready for use. Now you need to write some R code to read in and analyze all those file.

For your exercises




- You'll be asked to adapt your code from yesterday to run on the cluster for a much bigger ecological community size
- You'll need to collect species abundance data as before and average over a large number of parallel simulations.



For your exercises



- You'll be asked to adapt your code from yesterday to run on the cluster for a much bigger ecological community size
- You'll need to collect species abundance data as before and average over a large number of parallel simulations.
- Use a "burn in" period and check the species abundance distribution periodically. You should plot species richness against time and make a conservative judgment, but for neutral theory $4 * \text{metacommunity size}$ complete turnovers of the community is a good rule of thumb.



DO NOT ...

- Use the cluster without knowing memory and time requirements
- Run jobs on the login node
- Try to use cx2 or ax3 parts of the cluster
- Output data to the hard disk regularly
- Use the same random seed for your simulations
- Copy and paste your shell script
- Leave your results in \$TMPDIR
- Waste too much of your own time optimizing your code
- Run code on the cluster that hasn't been tested locally first

DO ...

- Use the cluster for jobs that take a long time locally.
- Optimize your code if there's going to be a huge benefit
- Run repeat readings and different parameters as separate jobs.
- Run jobs that take between 30 mins and 3 days to execute.
- Write your shell script on the cluster itself.
- Make your code output each result in a differently named file.
- Check periodically that all is well on the cluster
- Zip your output files to get them back to your PC.
- Be ambitious – you can do loads of great stuff with a cluster.