



High Performance Computing (HPC)

- Introduction to HPC and why it's useful
- How do you parallelize your code?
- The practice of running software on a cluster



Introduction to HPC and why it's useful

Gordon Moore



Born January 3, 1929 (age 84)
San Francisco, California, USA


Nationality American

Alma mater University of California, Berkeley;
California Institute of Technology

Occupation Chairman Emeritus, Intel
Corporation

Net worth ▲ \$4 billion USD (2011)^[1]

Jan 2015 update: \$6.7 billion



Introduction to HPC and why it's useful


- Embarrassingly parallel problems
 - Graphics
 - Simulations with multiple parameters
- Non embarrassingly parallel problems
 - Fluid dynamics
 - A lot of the tasks run by a single program




How do you parallelize your code? !

- `as.numeric(Sys.getenv("PBS_ARRAY_INDEX"))`
– Should be used in your code to give a simulation number

Before



Now




Using your PC

```
for ( i in 1 : 10 )
{
  do_simulation( i )
}
```

Using HPC


Shell script on
the cluster




How do you parallelize your code?

- `as.numeric(Sys.getenv("PBS_ARRAY_INDEX"))`
 - Should be used in your code to give a simulation number

Before




Now




How do you parallelize your code?


- `as.numeric(Sys.getenv("PBS_ARRAY_INDEX"))`
 - Should be used in your code to give a simulation number
- Pseudo random numbers
 - Given a certain random number seed, you get the same sequence of random numbers every time
 - Each simulation should have a different seed



Handling memory




Handling memory



Handling memory

- Save your results in memory and then write to disk at the end




Handling memory

- Save your results in memory and then write to disk at the end
- Output your code to a series of files
- Write local code to read in your series of files automatically



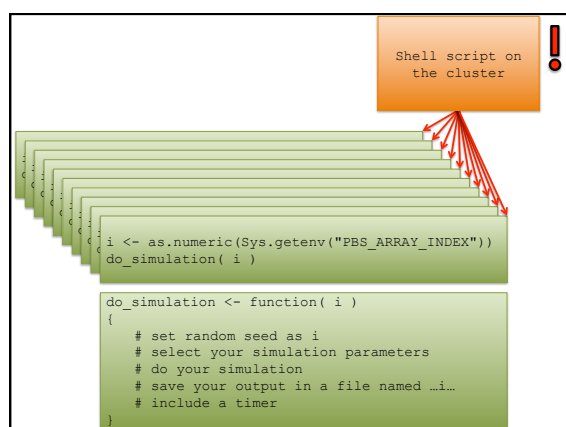
Handling memory

- Save your results in memory and then write to disk at the end
- Output your code to a series of files
- Write local code to read in your series of files automatically

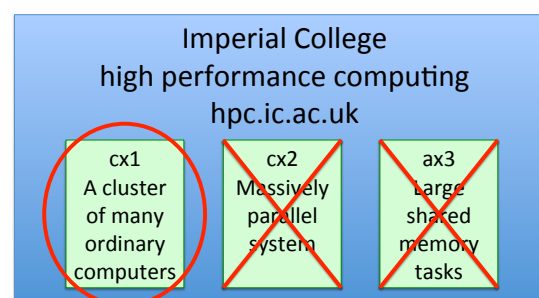


Handling memory

- Save your results in memory and then write to disk at the end
- Output your code to a series of files
- Write local code to read in your series of files automatically
- Build a timer into your code
- Test your code locally to know your memory and time requirements




The practice of running code on a cluster

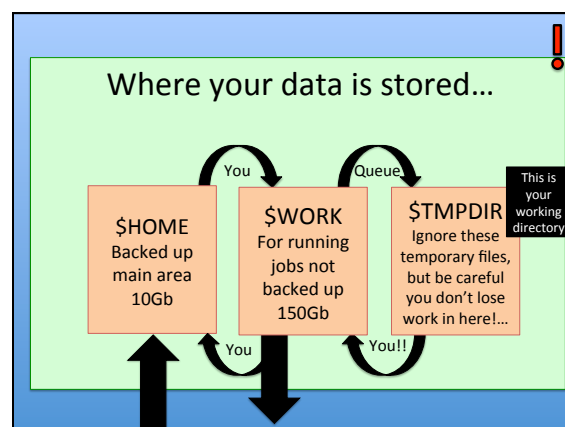
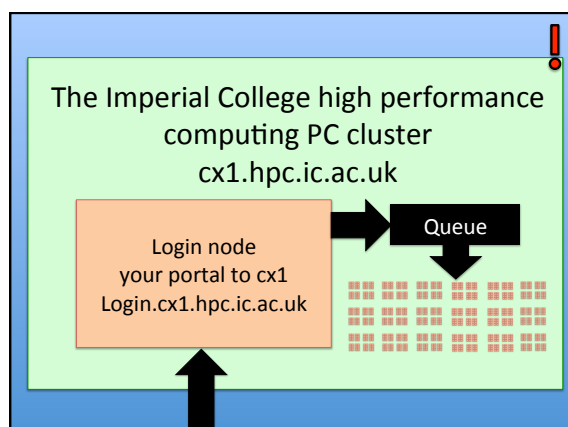


The Imperial College high performance computing PC cluster cx1.hpc.ic.ac.uk



The Imperial College high performance computing PC cluster cx1.hpc.ic.ac.uk





Step 1: get your code onto the cluster !

- Use sftp: from the directory of your code in a shell window type
 - sftp `username@login.cx1.hpc.ic.ac.uk`
 - You will be asked for your standard cluster password
 - put `filename.R`
 - exit
- Your software is now on your home directory of the cluster which is backed up regularly

Step 2: log into the cluster !

- Use ssh: from a shell window type
 - Ssh `-l username login.cx1.hpc.ic.ac.uk`
 - You will be asked for your standard cluster password
 - Now it's as though you were sitting with a shell open at the login node.
 - ls (will list the files in \$HOME)
 - mkdir `foldername` (make a new folder)
 - mv `filename $HOME/foldername` (move)
 - cd `foldername` (change directory)
 - cat `filename` (see your file to check it's contents)
 - cp `filename $WORK` (make a copy in \$WORK)
 - cd `$WORK` (change directory to \$WORK)
- Your software is now in two copies: in a folder in your home directory of the cluster and in the work area ready for running.

Step 3: make a file for your shell script !

- This is the list of instructions to be executed when you get to the front of the queue is written in shell script. It should be a .sh file
- **Never** run code on the login node – always write a shell script and wait in the queue.
- If you type `cat > filename.sh`
- You will then get the chance to type text (pressing enter for new lines) and the cat command will make the file containing the text that you typed.
- When you are finished typing the contents of your new file press control and D to complete the process.
- Type `cat filename.sh` to check that your file is correct before submitting it to the queue.

Step 3 continued: your shell script file !

```
#!/bin/bash
#PBS -l walltime=12:00:00
#PBS -l select=1:ncpus=1:mem=800mb
module load R/2.13.0
module load intel-suite
echo "R is about to run"
R --vanilla < $WORK/Rtest/ForwardsNTC_V5.R
mv datafilename* $WORK
echo "R has finished running"
# this is a comment at the end of the file
```

You can run Python code too –
just use different commands here

Step 4: submitting your job to the cluster!

- You are now in the \$WORK directory with your code and shell script both written.
- To submit your job type

```
qsub -J 1-32 filename.sh
```

```
qstat
```

(S changes from Q to B when running)
- If you want to delete a job

```
qstat
```

```
qdel job-id[]
```

(the [] is for array jobs only)
- qstat will give you a list of jobs and you would get the job-id from there.

Step 5: check that all is well

- Wait 5-10 minutes then check that nothing has gone wrong.
- qstat (is your job running still)
- ls (are output files as expected)
- cat filename.sh.ejob-id.index (are error files empty?)
- cat filename.sh.ojob-id.index (are standard output files as expected)
- qstat (is your job running still)
- exit (you're done for now come back later)

Step 6: Getting your results back

- qstat (is your job running still)
- cd \$WORK
- ls (output files as expected?)
- cat output filename (contents as expected?)
- cat filename.sh.ejob-id.index (error files empty?)
- cat filename.sh.ojob-id.index (standard output files as expected?)
- tar czvf filename.tgz *
- mv filename.tgz \$HOME

Step 6 continued: sftp to get results

- exit
- Use sftp: from a new directory on your own computer of where you want the results to be. Open a shell and type ...
 - sftp username@login.cx1.hpc.ic.ac.uk
 - You will be asked for your standard cluster password
 - get filename.tgz
 - exit
- Your results are now all on your own computer
 - tar xzvf filename.tgz
- Your results are now complete uncompressed and ready for use. Now you need to write some R code to read in and analyze all those file.